

Semi-Supervised Learning Task on Naver Fashion Dataset

Junseok Choi(20190665)
Team 7
w_choi@kaist.ac.kr

Minyoung Hwang(20170738)
Team 7
myhwang99@kaist.ac.kr

Junghyun Lee(20170500)
Team 7
jh_lee00@kaist.ac.kr

Abstract

In this paper, we propose a novel semi-supervised learning(SSL) framework for CS492(I) Naver Fashion Dataset challenge. First, we propose a framework in which multiple SSL algorithms are combined in an orthogonal manner. Also, by analyzing the given training dataset and exploiting some common characteristics among the images, we propose using attention-based pre-training. Through extensive empirical studies, we show that our approach gives a high performance, placing 2nd in the leaderboard.

1. Introduction

Semi-supervised learning (SSL), a learning framework where the training data is not labeled fully (usually, only a small percentage is labeled, and the rest is unlabeled), has received a lot of attention, lately. There has been two approaches to SSL methodology: regularization based method and pseudo-labeling based method. Regularization based method adds *unsupervised loss term* in order to account for the additional unlabeled data. Pseudo-labeling, first proposed in [6], is a wrapper method in which unlabeled datas are given *pseudo labels*.

As part of CS492(I), in this paper, we propose our solution to the SSL task with Naver Fashion Dataset. Unlike previous works, we propose a new framework in which various approaches are combined together to give better performance. (Specifically, the previously mentioned approaches fo SSL have been combined)

Our contributions can be summarized as follows:

- We’ve achieved 2nd place in CS492(I) leaderboard.
- By exploiting some of the common characteristics of the Naver Fashion Dataset, we propose using attention as part of our SSL-based approach.
- We show empirically that as long as one considers methods that introduces improvements in non-overlapping way (i.e. “orthogonal”), then combining them leads to much better accuracy.

1.1. Preliminaries

1.1.1 Naver Fashion Dataset/Resources

Naver Fashion Dataset has 58735 unlabeled images and 1060 training labeled images. Due to the nature of our task, we weren’t given any details of the test set. We note that even when compared to 1% of the ImageNet ILSVRC-2012[8], the size of the given training data was too small. We also had some constraints on the amount of resources provided by NSML[9].

1.1.2 Assumptions

In this subsection, we discuss some common underlying assumptions in SSL. These assumptions are necessary in the sense that if they are not satisfied, then additional unlabeled data leads to degradation in performance, implying breakdown of the SSL framework. Some (but not many) previous works have dealt with such phenomenon. (As stated in [15], this can be attributed to “publication bias”). For example, Cozman et al.[2] showed both theoretically and empirically that in some specific settings(mixture models) in which the assumptions do not hold, more unlabeled data may actually increase classification accuracy. (See also [3])

First is the manifold assumption, which asserts that the data actually lies on a very low-dimensional manifold, submerged in \mathbb{R}^d . Second is the smoothness assumption, which asserts if $x, x' \in X$ are close together in \mathcal{X} , then the corresponding labels y, y' should be the same. This is especially useful in SSL since if $x \in X_L$ and $x' \in X_U$ are close together, then there is a high probability (under this assumption) that x' should be labeled with the same label as x , which is already known. Lastly, we have the low-density assumption, which asserts that the decision boundary of a classifier must pass through low-density regions in \mathcal{X} . A simple argument[10] can be used to show that this assumption is actually a counterpart to the smoothness assumption.

2. Method

Our approach consists of two part; representation learning and fine tuning. Let us explain each part in detail,



Figure 1. The leftmost two images are selected from the Naver Fashion Dataset and the other were chosen from ImageNet.

then propose a method to combine all the mentioned algorithms. Along with that, we discuss how our method achieves superior performance compared to prior works, focusing specifically on how our method is very suitable for Naver Fashion Data set.

2.1. SimCLR

First proposed in [1], SimCLR is a simple framework for contrastive learning of visual representations. Suppose that we have a minibatch of size N .

A stochastic data augmentation: for each data x , two data transformations are performed to obtain x_1, x_2 . These two are called to be a *positive pair* since they have to be similar. Rest of the pairs are said to be *negative*. A combination of random crop and color distortion was used.

Contrastive loss function for a positive pair of samples is defined as follows:

$$l_{i,j} = -\log \frac{\exp(\text{sim}(z_i, z_j)/\tau)}{\sum_{k=1}^{2N} \mathbb{1}_{[k \neq i]} \exp(\text{sim}(z_i, z_k)/\tau)}$$

where $\text{sim}(u, v) = \frac{u^T v}{\|u\| \|v\|}$ is the cosine similarity. Using this loss, the base encoder $f(\cdot)$ and the projection head $g(\cdot)$ are updated.

2.2. Residual Attention Network

There is a clear distributive bias in Naver Fashion Dataset, which makes it very distinct compared to the commonly used datasets such as ImageNet([8]), as shown in Figure 1. Since ImageNet samples are photos taken from the nature without any editing, the backgrounds are continuous and not separated from the object. However, images of our dataset are manually edited so that the objects are usually separated from other features like the captions and the removed backgrounds and then centered. The bias can be exploited to improve performance by choosing a model that better identify the more important areas of the given image. Residual Attention Network([11]) (RAN) creates an attention mask that ranges from 0 to 1 and element-wise multiply to the feature map to decide

which part to pay more attention. The model then applies the mask on the previous feature map in a residual manner, so that the mask doesn't completely overwrite the useful features. This kind of ability makes it easier for the model to identify where the object is and pay less attention to the less important parts.

Optimization. We modify the network to exploit the GPU resources we have as much as possible. The model parameters and forward passing uses half-precision float values instead of single-precision. This makes it possible to ramp up the pre-train batch size to about twice of the case without it. Due to compatibility issues, the batch normalization layers and the optimizer use single precision computation. As SimCLR is heavily dependent on batch size[1], we thought that it outweighed the precision decrease coming from using half-precision.

2.3. Consistency Regularization

Similarly to SimCLR, we applied the idea of smoothness assumption for finetuning. Unsupervised data augmentation [12] is a method, where loss function consists of classification loss term over labeled data and consistency regularization term over unlabeled data. The brief flow of the work can be summarized as follows:

- For each given original input x , generate augmented pair \bar{x}, \tilde{x} .
- Predict distribution of labels for each augmented inputs $p_\theta(y|\bar{x})$, $p_\theta(y|\tilde{x})$ and optimize divergence $D(p_\theta(y|\bar{x}) \| p_\theta(y|\tilde{x}))$.

In this work, we differed the strength of the augmentations, considering the low accuracy of the model's prediction. Each weak and strong augmentation produces \bar{x}, \tilde{x} , respectively. During back propagation, the $p_\theta(y|\tilde{x})$ is enforced to mimic $p_\theta(y|\bar{x})$. Formally, the consistency loss term is defined as following:

$$I \left(\max_y p_\theta(y|\bar{x}) > \beta \right) CE(p_\theta(y|\bar{x}) \| p_\theta(y|\tilde{x})) \quad (1)$$

As proposed in prior work [12], we added condition indicating confidence or reliability of the prediction.

2.4. Label Propagation

Considering small number of labeled data set, the model may show too low accuracy to allocate pseudo-labels depending on itself. Instead, the model can learn good representation through unsupervised pre-task. It is reasonable to expect generating pseudo-labels based on inputs' representation vectors. The following explains a label propagation method introduced from a prior work [4].

According to manifold assumption and smoothness assumption mentioned in **Section 2**, we can predict distribution of labels over entire data space using labeled data points as anchors. In detail, we introduce adjacency matrix W , which encodes relative position between each point, and defines pseudo-labels Z of entire data as following:

$$Z = (I - \alpha W)^{-1} Y \quad (2)$$

Here, α is an arbitrary weight and Y is original labels, where labels for unlabeled data are initialized as *no-label* constant.

Practical procedure consists of three steps. To begin with, we make a nearest neighbor graph and conduct label propagation. Next, we decide weight of each label considering the reliability of propagated distributions. The final step is performing the normal classification task using pseudo-labels for unlabeled data. Detailed steps are same as the prior work [4].

2.5. Full Procedure.

To begin with, we train our self-attention model using SimCLR and move to finetuning. Model conducts classification task for finetuning. Finetuning procedure is divided into three sections. During the first section, model uses only labeled data. In the middle section, unlabeled data is inserted with pseudo labels resulted from label propagation. In last section, consistency regularization loss term is added with pairs of augmentation.

Advantages over previous approaches. Our approach is a combination of well-known previous methods; label propagation [4] and consistency regularization [12]. There are problems in using each one alone. In trivial, consistency regularization works well after model’s accuracy increases to some extent. Label propagation seems to be direct solution for the problem, but it is easy to converge to trivial solution, since augmentation on inputs for classification is weak. We expect that addition of consistency term may help the model to conduct classification task referring to general feature of inputs and escape such regions. We confirmed our belief through ablation test in **Section 4**.

3. Main results

All of our codes have been made publicly available in Github repository¹.

3.1. Effects of various components

3.1.1 Self-Attention

Since self-attention layer is simply added on top of the ResNet architecture, it was possible to compare the




Figure 2. Starting from the top left image, in clockwise: 1) a feature map before attention layer, 2) an input mask image, 3) an input shoe image, 4) a shoe image feature map after an attention layer, 5) a mask image feature map after an attention layer, 6) a feature map after attention layer.

Model	Top 1 Acc. (%)	Top 5 Acc. (%)
ResNet50	8.3	20.8
RAN56	42.3	22.6
ResNet101	4.5	8.7
RAN92	42.6	21.5

Table 1. Comparison of test set accuracy before label propagation using the best of each model.

performances where attention was used and where not. We believe the attention layer of our model is improving ResNet, which we verify quantitatively and qualitatively.

Quantitative Analysis. RAN architecture has the best test accuracy, and it seems that the additional attention layer enables the model to handle large number of parameters more efficiently, since simply using larger ResNet model with SimCLR didn’t result in a good accuracy. We hypothesize that this is due to the negative effects of small batch size on SimCLR (180 for ResNet101), where we’re forced to use when using a large model.

Qualitative Analysis. We examined the intermediate layer output of our model. From 1) and 6) of Figure 2 we can observe our model erasing the less important parts of feature maps where the object is at the middle and there are colored captions at both sides. In 2), 5) and 3), 4), the model focuses more on the objects and ignore the background, including its patterns.

3.1.2 Orthogonal combination of methods

This is shown through a series of experiments in which an incremental increase in performance is observed by adding on each method. The followings are details of our experiments, beginning with experiments applying consistency regularization and label propagation each alone, and our final proposal, combination of them.

¹https://github.com/cs492i-ELSA/cs492i_cv_repo_7




Figure 3. Accuracy of two models trained using consistency regularization and label propagation, respectively.

Consistency regularization. The training is performed for 20 epochs in total. Entire labeled and unlabeled data set is used as training set throughout the training. The loss term consists of classification loss for labeled data and consistency regularization term for unlabeled one. Random resized cropping to size 224×224 followed by random horizontal flipping is applied in common. The difference is addition of color distortion between them for strong augmentation. A batch consists of two sub-batches for 64 labeled and 64x30 unlabeled inputs. The initial learning rate l_0 is 0.02 and decays by 0.1 for every 6 epochs.

Label propagation. Whole procedure is divided into 2 stages. In first stage, classification task is conducted only for labeled data. This stage is up to 20 epochs. Before moving on to second stage, label propagation occurs and pseudo-labels are allocated to entire data including unlabeled ones. In second stage, model is trained by classification task over complete data set. Only weak augmentation in previous experiment is used. Batch consists of 2 labeled data and 32 unlabeled data. Learning rate maintained as 0.02 for 10 epochs and then decayed to 0.002 for left 14 epochs.

Consistency preserving label propagation. The brief procedure is introduced in Section 3. Hyper parameters are same as ones in previous experiment using label propagation. The one difference is that consistency regularization loss term is added to original classification loss since 10 epochs after label propagation. The term is for both labeled and unlabeled data.

Performance of each approaches. Figure 3.1.2 shows the progress of top-1 accuracy of training using consistency regularization and label propagation each alone. One using consistency regularization quickly converges to about 20%, which is easily achieved by using only labeled data. As we expected, label propagation have showed much better performance. Interestingly, Figure 3.1.2 and Table 3.1.2 proves that its performance gets even better after addition of consistency regularization.




Figure 4. Accuracy of model after addition of consistency loss against pure one using only label propagation.

Method	Best Top-1 Acc.	Best Score
CR	24.1	0.310
LP	25.7	0.328
LP + CR	26.8	0.329

Table 2. Comparison of three approaches.

3.1.3 Hyperparameter Tuning

We performed a learning rate sweep in $\{0.0005, 0.001, 0.002, 0.005, 0.01\}$ and selected the best optimizer among SGD, SGD+LARS([13, 7]), Adam([5]) and Yogi([14]). We also used both Exponential and Multi-step scheduler provided by PyTorch. Finally, we always used the largest batch possible for pre-training which was not big enough to add instability to the training and did a batch size sweep on the fine-tuning process, where we freeze the pre-trained model. Our best result was achieved using Adam, exponential scheduler and fine-tuning batch size 512. The performance degraded after size of 512, which we hypothesize is because of the large batch size and a lack of efficiency of the LARS optimizer when the pretrained model frozen and not updated.

4. Conclusion

In this work, we present a new approach to SSL by combining orthogonal procedures. We use an attention-based model that can efficiently exploit the dataset-specific characteristics. Then SimCLR was used for pre-training, and UDA and label propagation for fine-tuning, followed by hyperparameter tuning. We also perform an ablation study and verify our method makes a significant improvement over the baseline. By using our framework, we achieved 2nd place in the leaderboard.

References

- [1] Ting Chen, Simon Kornblith, Mohammad Norouzi, and Geoffrey Hinton. A simple framework for contrastive learning of visual representations. *arXiv preprint arXiv:2002.05709*, 2020.

- [2] Fabio Gagliardi Cozman, Ira Cohen, and Marcelo Cesar Cirelo. Semi-supervised learning of mixture models. In *Proceedings of the Twentieth International Conference on International Conference on Machine Learning*, ICML'03, page 99–106. AAAI Press, 2003.
- [3] David Elworthy. Does baum-welch re-estimation help taggers? In *Proceedings of the Fourth Conference on Applied Natural Language Processing*, ANLC '94, page 53–58, USA, 1994. Association for Computational Linguistics.
- [4] Ahmet Iscen, Giorgos Tolias, Yannis Avrithis, and Ondrej Chum. Label propagation for deep semi-supervised learning. In *IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2019, Long Beach, CA, USA, June 16-20, 2019*, pages 5070–5079. Computer Vision Foundation / IEEE, 2019.
- [5] Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. In Yoshua Bengio and Yann LeCun, editors, *3rd International Conference on Learning Representations, ICLR 2015, San Diego, CA, USA, May 7-9, 2015, Conference Track Proceedings*, 2015.
- [6] Dong-Hyun Lee. Pseudo-label: The simple and efficient semi-supervised learning method for deep neural networks. In *Workshop on Challenges in Representation Learning (ICML)*, volume 3, 2013.
- [7] Chunmyong Park, Heungsub Lee, Myungryong Jeong, Woonhyuk Baek, and Chiheon Kim. torchlars, A LARS implementation in PyTorch. <https://github.com/kakaobrain/torchlars>, 2019.
- [8] Olga Russakovsky, Jia Deng, Hao Su, Jonathan Krause, Sanjeev Satheesh, Sean Ma, Zhiheng Huang, Andrej Karpathy, Aditya Khosla, Michael Bernstein, Alexander C. Berg, and Li Fei-Fei. ImageNet Large Scale Visual Recognition Challenge. *International Journal of Computer Vision (IJCV)*, 115(3):211–252, 2015.
- [9] Nako Sung, Minkyu Kim, Hyunwoo Jo, Youngil Yang, Jingwoong Kim, Leonard Lausen, Youngkwan Kim, Gayoung Lee, Dong-Hyun Kwak, Jung-Woo Ha, and S. Kim. Nsm1: A machine learning platform that enables you to focus on your models. *ArXiv*, abs/1712.05902, 2017.
- [10] Jesper E. van Engelen and Holger H. Hoos. A survey on semi-supervised learning. *Machine Learning*, 109(2):373–440, 2020.
- [11] F. Wang, M. Jiang, C. Qian, S. Yang, C. Li, H. Zhang, X. Wang, and X. Tang. Residual attention network for image classification. In *2017 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 6450–6458, 2017.
- [12] Qizhe Xie, Zihang Dai, Eduard Hovy, Minh-Thang Luong, and Quoc V Le. Unsupervised data augmentation for consistency training. *arXiv preprint arXiv:1904.12848*, 2019.
- [13] Yang You, Igor Gitman, and Boris Ginsburg. Large batch training of convolutional networks, 2017.
- [14] Manzil Zaheer, Sashank Reddi, Devendra Sachan, Satyen Kale, and Sanjiv Kumar. Adaptive methods for nonconvex optimization. In S. Bengio, H. Wallach, H. Larochelle, K. Grauman, N. Cesa-Bianchi, and R. Garnett, editors, *Advances in Neural Information Processing Systems 31*, pages 9793–9803. Curran Associates, Inc., 2018.
- [15] Xiaojin Zhu. Semi-supervised learning literature survey. Technical Report 1530, Computer Sciences, University of Wisconsin-Madison, 2005.